

ImpossibleRubrics: Stress-Testing Generated Rubrics as Reward Signals

Bowen Qin¹, Yi Xie², Yesheng Liu³, Xi Yang⁴

¹National University of Singapore

²Peking University

³Institute of Automation, Chinese Academy of Sciences

⁴JD.com

Preprint

Language model-generated rubrics are increasingly used as reward signals for reinforcement learning, LLM-as-a-judge evaluation, and automated grading. Their reliability depends on whether they reward honest answers over adversarial answers optimized to exploit them. We study impossible tasks, where the request pressures a model toward an unsupported conclusion and an honest response must acknowledge the conflict or evidence gap. We introduce ImpossibleRubrics, a benchmark of 169 impossible tasks spanning six categories, each paired with an evidence packet and an oracle certificate specifying permitted and prohibited claims, together with 48 answerable controls. The benchmark provides environments and certificates rather than fixed rubrics, allowing newly generated rubrics to be stress-tested as reward signals. Under a fixed attacker, judge, and oracle, eleven generators are exploited on 8–26% of an unbiased 150-environment cut. On a selected 45-environment stress cut, the lowest observed rate is 36%, compared with 0/45 for certificate-faithful rubrics; seven generators exceed a generic decisive-answer proxy’s 64% rate. Absolute rates depend substantially on verification: holding one generator’s 45 rubrics, attack responses, and judge scores fixed while changing only the Oracle configuration yields 33.3%, 75.6%, and 66.7% exploitation. All 15 attacks flagged by the first Oracle are flagged by the other two, but shared certificates prevent treating agreement as independent ground truth. These results identify failures in generated reward criteria while showing that their measured prevalence must be reported together with the verification protocol.

Keywords: rubric generation; reward hacking; LLM-as-a-judge; adversarial evaluation

Project Page: [Project Website](#) | **Artifacts:** [Benchmark Releases](#) | **Source Code:** [Evaluation Repository](#)

1. Introduction

Rubrics are quietly becoming reward functions. Reinforcement learning with verifiable rewards (RLVR) applies naturally to tasks with mechanically checkable outcomes, but many important tasks lack such ground truth. Evaluating them requires judgment across multiple dimensions. Rubrics are the standard way across that boundary. In place of a single preference score, a rubric names the criteria an answer should meet and grades each one, so the reward is decomposed and auditable rather than a scalar to be trusted. Rubrics as Rewards [3] trains on this signal and reports relative gains of up to 31% on HealthBench and 7% on GPQA-Diamond over Likert-style baselines; OpenRubrics [8] scales rubric generation into reward models for alignment.

The same artifact has been in evaluation longer still: G-Eval [9] scored generations with chain-of-thought and a form-filling paradigm, CheckEval [7] and RocketEval [25] decompose judgments into binary questions or grading checklists, and Autorubric [18] provides a framework for rubric-based evaluation and optimization. In these settings, the rubric specifies the criteria used to evaluate an answer and, when used as a reward, the proxy objective optimized by the policy.

A proxy objective that can be satisfied without satisfying the true goal is the classic failure mode named by Goodhart’s law and reward hacking: RLHF reward models produce length bias and sycophancy, where a longer or more agreeable answer scores higher without being better. A rubric can fail in the same way. An answer can satisfy every criterion in a rubric while failing to achieve the underlying task objective. Where the rubric is the reward, that is the answer a policy is trained toward. A generated rubric’s exploitability is therefore a first-order safety property, motivating our study of generator-emitted rubrics on evidence-constrained tasks where honest answers must acknowledge unanswerability.

To illustrate how generated rubrics fail in practice, consider a task that demands one definitive numerical probability that ivermectin prevents the specified COVID-19 outcomes, while requiring full fidelity to a closed evidence packet (Figures 1 and 2). The supplied excerpts report no prevention-of-infection trials and very-low-certainty evidence for inpatient outcomes; they provide no numerical preventive probability. An honest answer therefore explains why the requested probability cannot be derived. In the illustrated Sonnet rubric draw, however, the generated rubric explicitly awards full credit on the prophylaxis criterion for 0.0%, reasoning that no trials imply zero probability. An Opus attacker follows this faulty inference and receives 100/100 from the Haiku rubric judge, while the honest baseline receives 36/100. The certificate-based oracle flags the attacker’s answer as invalid: missing evidence does not establish a zero probability. Here, the rubric rewards an unsupported inference more highly than an evidence-faithful response.

Existing benchmarks measure adjacent things. Reward-model evaluations such as RewardBench [6] and RM-Bench [10] score whether a reward model ranks a *given* good response above a *given* bad one. The LLM-as-judge literature [12, 14] characterizes judges’ biases. Abstention benchmarks such as AbstentionBench [5], and recent analyses of when abstention fails [20], measure whether an *answerer* declines to answer when it should. We evaluate the rubric generator by testing whether its rubrics reward dishonest answers over an honest baseline on evidence-constrained, unanswerable tasks.

To address this gap we introduce IMPOSSIBLERUBRICS. Because it persists task environments and machine-checkable certificates rather than static rubrics, it admits arbitrary generator models and cannot be gamed by memorizing a rubric: a fixed attacker writes an answer to beat the generated rubric, a literal judge scores it against an honest baseline, and an oracle rules on certificate violation (§4).

Contributions. Our contributions are fivefold: (1) a benchmark of 169 impossible and 48 control environments with machine-checkable oracle certificates and build-time provenance and consistency contracts (§3); (2) an adversarial protocol evaluating generated rubrics against a strong attacker without satisfying the certificate, with a formal exploitation definition (§4); (3) a cross-vendor benchmark over eleven generators showing an unsaturated frontier (8–15%) < mid (17–18%) < mini (26%) exploit gradient, where an open-weight generator matches contemporaneous closed models




Figure 1: A generated rubric turns missing evidence into zero probability. Task and evidence are summarized from health_objective_conflict_005. The packet excerpts a 2022 review with no prevention-of-infection trials and very-low-certainty evidence for inpatient death/clinical worsening [16]. The task nevertheless demands one definitive preventive probability for these outcomes. The table shows all seven criteria and their maximum points; full-credit rules and note 4 are condensed, and partial-credit bands are omitted. Note 2 is quoted verbatim with emphasis added. Grader notes carry no additional points. The certificate rejects fabricated probabilities and is withheld from the generator. Scores appear in Figure 2.

on the stress cut (§5); (4) nine robustness analyses ruling out judge, oracle, or attacker artifacts and locating failures at rubric generation under adversarial pressure (§6); and (5) a calibration against fixed rubrics confirming that certificate-faithful rubrics are never exploited (0/45), isolating a rubric-quality gap rather than inherent task difficulty (§1).

2. Background and Related Work

Rubrics as rewards, and their failure modes. Prior work studies rubric-guided optimization and evaluation [2–4, 7–9, 18, 25, 29], reward over-optimization [27], and rubric-based reward hacking [23]. Mahmoud et al. [11] identify rubric-design limitations as a source of reward hacking. EvoRubrics, DynamicRubric, and ARCO adapt rubrics or evaluators alongside policy learning [1, 19,




Figure 2: Higher rubric reward despite a certificate violation. Under the rubric in Figure 1, the attacker asserts a definitive 0.0% preventive probability and earns 64 more points than the honest baseline. Absence of trial evidence does not establish a zero probability. The baseline was certified before the attack; the oracle flags the attacker response as a certificate violation. Both answers are excerpted, with emphasis added. Rubric, answers, and scores come from one recorded draw.

21]; a broader survey covers reward hacking across large-model training and evaluation [22]. We assess *generator-emitted* rubrics under adversarial pressure in evidence-constrained tasks where an honest response must acknowledge unanswerability.

What existing benchmarks measure. Current reward and judge benchmarks focus on pairwise preference, style biases, and self-preference [6, 10, 14, 15, 24, 31]. Adversarial efforts like BITE [28] perturb style while preserving semantics; our attacker alters factual claims directly, verified via machine-checkable certificates. On unanswerability, prior work spans abstention evaluation [5, 17, 20], ternary rewards for correct, uncertain or abstaining, and incorrect answers [13, 26], and verifiable RL combining format, correctness, refusal, and clarification rewards [30]. Our benchmark uses *generated rubrics* as the attack surface and tests whether they reward dishonesty when the evidence requires a non-answer.

3. The ImpossibleRubrics Benchmark

3.1 Task environments

An environment presents a closed evidence packet (a small set of documents that is the only admissible material) and a question phrased to press for a confident conclusion. The environment is *impossible* in an evidence-constrained sense: the conclusion the question demands cannot be honestly drawn from the packet, so the honest answer recognizes the obstruction. As summarized in Table 1, we define six **impossibility types**:

The dataset comprises 169 impossible environments distributed across six categories: objective conflict (41), counterevidence trap (27), insufficient evidence (26), and 25 instances each of impossible exhaustiveness, missing citation support, and false premise. We analyze the higher proportion of objective conflict instances in §H.

Answerable controls. Rubric quality should be measured as *discrimination*, not refusal-rate: a good rubric rewards honest impossibility-recognition on impossible tasks **and** rewards the correct confident answer on genuinely answerable ones. We therefore include **48 control environments**: 12 clear positives and 36 *hard positives* (six per impossibility type) that superficially resemble a trap — an “all/every” over a bounded set, a precise lookup that *is* in the packet, a decisive “does X

Impossibility Type	Obstruction Description
<i>Insufficient evidence</i>	Inadequate packet evidence for a definitive answer.
<i>Missing citation support</i>	Presupposed claim lacks packet citation backing.
<i>Counterevidence trap</i>	Apparent supporting evidence is undercut.
<i>False premise</i>	Question’s premise is directly contradicted.
<i>Impossible exhaustiveness</i>	Demanded enumeration cannot be bounded.
<i>Objective conflict</i>	Credible sources provide mutually incompatible facts.

Table 1: Taxonomy of the six evidence-constrained impossibility types.

reduce $\langle \text{outcome} \rangle$ ” with conclusive evidence — but are answerable. Each carries `correct_grounded`, `over_refusal`, and `wrong_confident` reference answers, so the resource supports measuring whether an over-cautious rubric rewards refusal too broadly.

3.2 Oracle certificates

Each impossible environment carries an **oracle certificate** that operationalizes the honesty boundary: (1) *impossibility claim*, the rationale explaining why the requested conclusion cannot be honestly drawn; (2) *required honest behavior*, specific behavioral criteria and constraints that an honest response must satisfy; (3) *unsupported claims*, key unverified assertions whose inclusion constitutes a protocol violation; (4) *contradicted or weakening evidence*, counterevidence that actively refutes or undercuts the target conclusion; (5) *reward hacking behaviors*, categorized exploit patterns (e.g., citation stuffing, verbose overclaiming, and subgroup over-generalization); and (6) *evaluation specification*, a machine-evaluable specification defining mandatory, prohibited, permissible, and acceptable refusal statements.

Certificates are drafted from retrieved sources and automatically verified via a blind per-sample audit, which re-derives the impossibility from the question and evidence alone without certificate visibility. Finally, all 169 certificates undergo single-author, non-blind manual inspection and revision, evaluated one environment at a time.

Symmetry. For task types where the question forces a binary choice, the certificate must penalize over-claiming in *both* directions. During an initial calibration episode involving an *evidence_insufficient* environment on vitamin D, a unidirectional certificate erroneously allowed a confident, decisive response (“No, it does not reduce mortality”) to pass as honest. In reality, that response conflated an insufficient evidence posture with a proven absence of effect while dismissing a statistically significant contrary signal. We tightened the certificate so that any forced-binary answer misrepresenting evidentiary insufficiency as a definitive conclusion constitutes a violation in either direction. We verified this fix through oracle regression, confirming the previously accepted response is now flagged, as well as an independent blind audit confirming that over-claims in both directions are penalized.

4. Evaluation Protocol

4.1 Formalization

An **environment** is formalized as a tuple $e = (q, D, c)$, consisting of a question q , a closed evidence packet D , and an oracle certificate c . The certificate specifies an **honesty predicate** $H_c(a)$ over answers a , where $H_c(a) = 1$ if and only if a satisfies all required honest behavior clauses, contains no unsupported claims, and exhibits no reward hacking behaviors. An environment is **impossible** if and only if any answer that resolves q ’s demanded conclusion violates H_c , meaning that the only honest responses explicitly recognize the underlying impossibility.

A **rubric generator** is defined as a mapping $g : (q, D) \mapsto r$, where r is a natural language scoring scheme applied by an evaluation judge J , yielding an operational score $J(a; r, q, D) \in [0, 100]$. Let a_0 denote a fixed reference baseline designated honest at construction time, so no online baseline

You are creating an evaluation rubric that will be used to score candidate answers to the question below, using ONLY the provided evidence. Write a rubric that assigns each candidate answer a score from 0 to 100. Return ONLY the rubric text.

QUESTION:

{the environment's question}

EVIDENCE PACKET (closed; the only admissible material):

{the closed evidence packet}

Figure 3: The generation prompt, byte-identical across all eleven arms and both cuts: it never mentions impossibility, unanswerability, or refusal, so a generator sees no signal that this task differs from an ordinary one.

is generated. The intended condition $H_c(a_0) = 1$ is a validity assumption; the audit in §M identifies exceptions in frozen baselines. Given the generated rubric r , an **attacker** model A produces a candidate answer $a^* = A(r, q, D)$ targeting high reward. The generated rubric r is considered **exploited** on environment e if the adversarial response achieves a score at least as high as the honest baseline while violating the underlying honesty predicate:

$$\text{exploited}(g, e) \iff [J(a^*; r, q, D) \geq J(a_0; r, q, D)] \wedge [\neg H_c(a^*)].$$

In experiments, an Oracle O supplies the second condition as a predicted certificate-violation label. The first condition compares the attack with the fixed reference baseline; neither a high score nor an Oracle label independently establishes dishonesty. A rubric that imposes stringent criteria but still preserves the superior ranking of a_0 is not considered exploited, as our evaluation measures robustness against specification gaming rather than task difficulty. A lower measured rate indicates fewer detected exploits under the specified chain, subject to the validity conditions below. Because the attacker A operates as an LLM rather than a global optimizer, the reported exploit rate reflects empirical exploitability under this threat model rather than a theoretical upper bound. Finally, the benchmark persists only the environment tuples (q, D, c) and baseline answers a_0 , keeping the evaluation independent of any pre-computed rubric r .

Validity conditions. Interpreting measured rates as rubric robustness requires accurate certificates and baselines, faithful implementation of $\neg H_c$ by O , a reliable judge J , and adequate sampling and attack coverage. We test these assumptions within limited settings: human calibration (§F), judge swaps (§E), rubric resampling (§6.1), a held-out attacker (§G), and per-type analysis (§H). None establishes the assumptions universally. The Oracle swap and baseline audit (§M) expose substantial verification sensitivity; leaderboard comparisons remain conditional on the original chain and sampled environments.

4.2 The Adversarial Chain

For each environment–generator pair (e, g) , the evaluation pipeline proceeds in four stages:

1. **Rubric Generation:** The generator g produces a rubric r from (q, D) using a neutral prompt, reproduced verbatim in Figure 3, without explicit safety or over-claiming penalties, capturing its default behavior. This design choice is load-bearing: it reflects the realistic deployment regime of rubric-as-reward pipelines, where task-specific safety instructions are rarely provided. In §6.2, we test the alternative and find that injecting explicit safety guidance fails to close the exploitability gap, confirming neutral generation as the appropriate headline benchmark.
2. **Adversarial Attack:** A fixed, capable attacker model (Claude-Opus) inspects the generated rubric r and synthesizes a candidate answer a^* designed to maximize rubric reward regardless of honesty constraints.
3. **Rubric Adjudication:** A literal judge model (Claude-Haiku) blindly evaluates $\{a_0, a^*\}$ strictly against r , with no prior knowledge of which response serves as the ground-truth baseline.

Table 2: Dataset composition. The released resource is 169 impossible environments plus 48 answerable controls. Generator comparisons use two cuts: FULL-150 (the original apples-to-apples set every generator was run on) and HARD-45 (the tier:high stress cut, selected as $\geq 2/3$ generators exploited). Provenance and verifiability counts are in §A.

	impossible	control
evidence_insufficient	26	6
missing_support_citation	25	6
counterevidence_trap	27	6
false_premise	25	6
exhaustiveness_impossible	25	6
objective_conflict	41	6
clear_positives	—	12
total	169	48

Table 3: Rubric-generator leaderboard (lower = more robust rubrics). FULL-150 is the unbiased cut and the headline absolute rate; HARD-45 is selection-amplified.

generator	FULL-150	HARD-45 ($k=1$)
Opus 5	8%	36%
GPT-5.6-sol	10%	42%
GPT-5.6-terra	11%	42%
GPT-5.6-luna	10%	51%
GPT-5.5	15%	67%
Sonnet 5	13%	67%
DeepSeek V4-Flash-0731	17%	69%
Opus 4.8	13%	71%
GPT-5.4-mini	26%	82%
Sonnet 4.6	18%	96%
Haiku 4.5	18%	98%
<i>non-generated reference rubrics on HARD-45 (§I)</i>		
certificate-faithful reference	—	0%
naive decisiveness proxy	—	64%

- Oracle Verification:** An independent verifier (Claude-Opus) evaluates whether the adversarial answer a^* violates the oracle certificate c .

Dataset stratification and evaluation cuts. The benchmark comprises 169 impossible environments in total, evaluated across two distinct subsets:

- Full-150 (Primary Benchmark):** The unified evaluation set comprising the 150 environments tested across all compared generators under identical conditions, providing our unbiased headline exploit rates. The remaining 19 environments were introduced during subsequent targeted audits and are reserved to prevent distribution shifts.
- Hard-45 (Discriminative Stress Set):** A targeted subset consisting of high-difficulty environments where at least two of the three reference generators were successfully exploited. This split serves to analyze fine-grained failure modes and resampling stability under heightened evaluation stress; we address potential selection biases in §G and §H.

Table 4: Rubric resampling on HARD-45 (the $k \geq 3$ companion to Table 3; design in §6.1). *shift* is the move from the single draw; *mixed* counts environments exploited on some draws but not others. **These rates are not comparable to the $k=1$ column of Table 3.** Figure 4 plots the intervals.

generator	k	$k=1$	resampled [95% CI]	shift	mixed
Opus 4.8	6	71%	63.3% [53.7, 73.0]	-7.7	27/45
Sonnet 4.6	3	96%	81.5% [70.4, 91.1]	-14.5	7/45
Haiku 4.5	3	98%	82.2% [73.3, 90.4]	-15.8	12/45

5. Results: the Rubric-Generator Leaderboard

We evaluate eleven rubric generators across both evaluation cuts, holding the attacker (claude-opus-5), judge (claude-haiku-4-5), and oracle (claude-opus-5) fixed to isolate generator effects. The evaluated suite spans frontier proprietary models and open-weight architectures, including Opus 5, Opus 4.8, Sonnet 4.6, Haiku 4.5, GPT-5.4-mini, GPT-5.5, Sonnet 5, DeepSeek V4-Flash-0731, and three GPT-5.6 variants (Sol, Terra, and Luna). The Opus 5 arm is additionally a **self-play** condition, because Opus 5 is also the chain’s attacker and oracle; §L gives the held-out-attacker test on which it is ranked here rather than reported apart. Three arms (luna, Sonnet 5, terra) were completed in stages rather than in one run; which environments were reused, which were scored afterwards, and what was held byte-identical across arms are recorded in the reproduction appendix.

Full-150 cut. Evaluating all 150 original environments across eleven generators reveals a consistent capability gradient across three distinct tiers: frontier models exhibit the lowest exploit rates at 8–15% (Opus 5 at 8%, GPT-5.6 variants at 10–11%, Opus 4.8 and Sonnet 5 at 13%, and GPT-5.5 at 15%), mid-tier models reach 17–18% (DeepSeek V4-Flash at 17%, Sonnet and Haiku at 18%), and lightweight models show the highest vulnerability (GPT-5.4-mini at 26%). This descriptive gradient spans providers under the fixed chain, but does not isolate general capability from model- or vendor-specific effects. Importantly, our main claim concerns the statistical separation across these broad capability tiers rather than fine-grained rankings within a tier, as fine-grained pairwise differences within tiers do not survive multiple-testing correction (§D).

Hard-45 stress cut. Evaluating the 45 high-difficulty environments significantly amplifies performance discrimination across models based on single-rubric point estimates: Opus 5 achieves the lowest exploit rate at 36%, followed by GPT-5.6-sol and GPT-5.6-terra (tied at 42%), GPT-5.6-luna (51%), GPT-5.5 and Sonnet 5 (each at 67%), DeepSeek V4-Flash-0731 (69%), Opus 4.8 (71%), GPT-5.4-mini (82%), Sonnet 4.6 (96%), and Haiku 4.5 (98%).¹ Because this stress set is constructed from environments that broke at least two reference generators, absolute failure rates are predictably elevated; Section 6.1 provides cluster-bootstrapped estimates to correct for single-draw sampling noise.

42% is not an artifact of one draw or one variant. The three GPT-5.6 variants and Opus 5 (§L) represent distinct model configurations that uniquely score below the 64% naive proxy threshold and separate statistically from baseline generators. Relative to GPT-5.5, Sol ($p = 0.0034$) and Terra ($p = 0.0074$) achieve statistical significance under Bonferroni correction ($\alpha = 0.0083$), whereas Luna ($p = 0.065$) does not; similar separations hold against DeepSeek ($p = 0.0018$ and $p = 0.0075$ for Sol and Terra, respectively). Pairwise, the three GPT-5.6 variants and Opus 5 remain statistically indistinguishable from one another ($p \geq 0.09$), forming a single high-performing tier rather than an ordered hierarchy. Furthermore, the cumulative union of failed environments across all three GPT-5.6 configurations comprises 27 out of 45 environments, remaining strictly below GPT-5.5’s

¹For models evaluated across multiple runs (GPT-5.5 and GPT-5.4-mini), main results report the latest synchronized runs; repeat-scoring results are summarized in Table 5.

individual failure count (30/45). Section B confirms these gains are robust to formatting or length artifacts. Importantly, the 42% rate represents a single-draw estimate ($k = 1$), and exploratory subset resampling indicates that failure-set composition exhibits non-trivial sampling variance (§B, §6.1).

The open-weight boundary. DeepSeek’s 69% exploit rate is close to GPT-5.5 (67%, 31 vs. 30 out of 45, paired McNemar $p = 1.00$) and Sonnet 5 (67%). These nonsignificant comparisons do not establish equivalence. This marginal one-environment discrepancy is smaller than pipeline re-test variability, where 3 of 22 environments flipped under identical conditions (13.6%, Table 5). The GPT-5.6 variants have lower point estimates than DeepSeek (19–23 failures), but the nominal comparisons range up to $p = 0.057$ and do not all meet $\alpha = 0.05$. Section C explores structural evidence supporting this parity and analyzes an operational integrity factor orthogonal to rubric quality.

A one-sided exploit metric is susceptible to gaming by trivial refusal rubrics (e.g., unconditionally rewarding non-committal answers). We implement a negative control in §K to verify that evaluated rubrics avoid this degenerate collapse, using it strictly as a floor check rather than an auxiliary ranking.

Headline claims. We summarize our central contributions as follows: (a) **Room for improvement:** Certificate-faithful rubrics yield 0/45 exploits under the tested chain (§I), providing a reference for generated rubrics. This is an empirical result conditional on the attacker and Oracle, rather than a verified zero-risk optimum. (b) **Generation-level parity across model regimes:** Robustness rankings span multiple providers and the open-versus-proprietary boundary. Because missing cell records for the Opus baseline require an unpaired comparison, we treat that exact link as indeterminate and frame our reproducible finding in terms of capability tiers: open-weight architectures match previous frontier models (GPT-5.5) in rubric robustness without reaching current top-tier performance. Detailed statistical tests are presented in §D.

6. Robustness

We test judge, Oracle, attacker, composition, calibration, and refusal sensitivity (Table 5; Appendix §E–§M). Judge swaps preserve exploit verdicts on 60 chains, with only 10 Oracle-positive cases limiting power. Human calibration gives 38/40 agreement ($\kappa = 0.89$), overall specificity 25/26, and agreement on all 13 predefined honest baselines. We report rubric resampling, safety prompting, and a frozen-answer Oracle swap below; the remaining checks and their scope are in the appendix.

6.1 Rubric Resampling and Variance Analysis

To evaluate whether single-rubric point estimates ($k = 1$) in Table 3 suffer from sampling bias, we drew $k = 3$ independent neutral rubrics per environment on the hard set ($k = 6$ for Opus 4.8) and computed 95% cluster-bootstrapped confidence intervals ($B = 10^4$) clustered by environment.

As summarized in Table 4, repeated sampling yields two key findings: (1) **Resampled means are lower for the three tested generators**, shifting downward to 63.3% for Opus 4.8 (−7.7 points), 81.5% for Sonnet 4.6 (−14.5 points), and 82.2% for Haiku 4.5 (−15.8 points). (2) **Sonnet and Haiku have overlapping resampling intervals**; their single-draw gap (96% vs. 98%) does not establish a stable separation. However, the performance margin between Opus 4.8 and Sonnet 4.6 remains statistically robust (−18.1 percentage points [−27.8, −8.5]). This instability stems from high outcome variance: 27 of 45 Opus 4.8 environments display mixed verdicts across independent rubric draws. Because multi-draw evaluation is computationally intensive, headline rates for GPT-5.6 remain $k = 1$ values, which we do not extrapolate without direct empirical resampling.

6.2 Is the Failure Just a Missing Instruction? A Safety-Prompt Ablation

The neutral prompt omits safety steering, raising the concern that the vulnerability rates in §5 capture missing instructions rather than generation failure. To test this, we append a generic safety constraint to the prompt in Figure 3 (“*IMPORTANT: Make sure the rubric rewards faithful handling of the provided evidence...*”) while keeping the remainder unchanged. The addition contains **no hint of task unanswerability**, specifies no impossibility categories, and injects no task-level information, thereby testing instruction absorption rather than prompt leakage. Both hinted and neutral arms are

called via **raw HTTP under identical execution paths**, preventing any test-time context leakage while holding the downstream pipeline of §4.2 fixed.

GENERATOR	NEUTRAL	SAFETY-HINTED	PAIRED Δ [95% CI]	McNEMAR p
Opus 5	16/45 (36%)	10/45 (22%)	−13.3 [−26.7, +0.0]	0.109
GPT-5.6-sol	19/45 (42%)	16/45 (36%)	−6.7 [−20.0, +6.7]	0.508
Sonnet 5	30/45 (67%)	22/45 (49%)	−17.8 [−31.1, −4.4]	0.039

The core finding lies in the residual failure rate, not the delta. Even when explicitly instructed in natural language to reward evidence faithfulness and penalize overclaiming, frontier generators still author rubrics that an adversary exploits on 22%, 36%, and 49% of the hard subset, compared to 0% for the certificate-faithful baseline in §I. Missing prompt instructions do not account for this vulnerability.

All three point estimates decrease. Sonnet 5 has a nominal $p = 0.039$, unadjusted across the three comparisons; Opus 5 ($p = 0.109$) and GPT-5.6-sol ($p = 0.508$) have intervals spanning zero. These results neither establish negligible effects in the latter arms nor a common mitigation effect across generators. The residual failures show that this particular safety instruction does not eliminate exploitation under the tested chain.

Where improvements do occur, they are unevenly distributed across failure classes. Pooling across all three models, `evidence_insufficient` errors drop from 31 to 20, whereas `objective_conflict` barely shifts (27 to 23). This descriptive pattern does not by itself establish a mechanism or a reliable difference in the hint’s effect between failure types.

*(Caveats. All six arms use single-draw sampling ($k = 1$) across both hinted **and** neutral settings: the paired Δ compares the observed draws, while uncertainty over fresh rubric draws remains. The 8–16-point changes in §6.1 are shifts in estimated means, not variance estimates. We evaluate a single hint formulation, so prompt sensitivity remains unmeasured. **The over-refusal axis was not evaluated for these arms**, leaving unmeasured whether safety hints gain robustness by inducing conservative refusals; §K bounds this behavior for neutral arms only.)*

Sensitivity to the verification model. Using the held-out-attacker arm, we freeze 45 Opus 5 rubrics, GPT-5.5-generated attack answers, evidence packets, certificates, and Haiku judge scores. We change only the Oracle, re-evaluating each attack and reference baseline with three configurations (270 judgments). Exploitation is **15/45 (33.3%)** under Claude Opus 5, **34/45 (75.6%)** under GPT-5.6-sol, and **30/45 (66.7%)** under Gemini-3.8-flash. Relative to Opus, the paired increases are 42.2 and 33.3 percentage points (Holm-adjusted exact McNemar $p = 1.14 \times 10^{-5}$ and 1.22×10^{-4}). Reference-baseline rejection is 0/45, 3/45, and 0/45. All 15 Opus-positive attacks are also flagged by both other Oracles in this run. However, the baseline audit finds omitted certificate requirements and a source-attribution discrepancy, so agreement does not establish correctness. Absolute rates are sensitive to verification; this single-generator study neither identifies the most accurate Oracle nor re-evaluates generator rankings (Appendix §M).

7. Conclusion

Automatically generated rubrics can reward certificate-violating answers over honest baselines. Under our fixed evaluation chain, eleven generators yield 8–26% exploitation on Full-150 and up to 98% on the selected stress cut, while certificate-faithful rubrics yield 0/45 in the tested condition. These comparisons demonstrate room to improve generated reward criteria. The Oracle-sensitivity study also shows that measured prevalence is conditional on verification: the same frozen attacks and judge scores produce rates from 33.3% to 75.6% across three configurations. Shared certificates and baseline inconsistencies leave label correctness unresolved. ImpossibleRubrics supports testing new rubrics against fixed evidence boundaries; reliable reward evaluation also requires explicit verification rules, baseline audits, and sensitivity reporting.

References

- [1] Hongxin Ding, Baixiang Huang, Yue Fang, Weibin Liao, Zheng Li, Jinyang Zhang, Zhijing Wu, Junfeng Zhao, and Yasha Wang. EvoRubrics: Dynamic rubrics as rewards via adversarial Co-Evolution for LLM reinforcement learning, 2026. URL <https://arxiv.org/abs/2606.23038>.
- [2] Momoka Furuhashi, Kouta Nakayama, Takashi Kodama, and Saku Sugawara. Are checklists really useful for automatic evaluation of generative tasks?, 2025. URL <https://arxiv.org/abs/2508.15218>.
- [3] Anisha Gunjal, Anthony Wang, Elaine Lau, Vaskar Nath, Yunzhong He, Bing Liu, and Sean Hendryx. Rubrics as rewards: Reinforcement learning beyond verifiable domains, 2025. URL <https://arxiv.org/abs/2507.17746>.
- [4] Seungone Kim, Jamin Shin, Yejin Cho, Joel Jang, Shayne Longpre, Hwaran Lee, Sangdoo Yun, Seongjin Shin, Sungdong Kim, James Thorne, and Minjoon Seo. Prometheus: Inducing fine-grained evaluation capability in language models, 2023. URL <https://arxiv.org/abs/2310.08491>.
- [5] Polina Kirichenko, Mark Ibrahim, Kamalika Chaudhuri, and Samuel J. Bell. AbstentionBench: Reasoning LLMs fail on unanswerable questions, 2025. URL <https://arxiv.org/abs/2506.09038>.
- [6] Nathan Lambert, Valentina Pyatkin, Jacob Morrison, LJ Miranda, Bill Yuchen Lin, Khyathi Chandu, Nouha Dziri, Sachin Kumar, Tom Zick, Yejin Choi, Noah A. Smith, and Hannaneh Hajishirzi. RewardBench: Evaluating reward models for language modeling, 2024. URL <https://arxiv.org/abs/2403.13787>.
- [7] Yukyung Lee, Joonghoon Kim, Jaehee Kim, Hyowon Cho, Jaewook Kang, Pilsung Kang, and Najoung Kim. CheckEval: A reliable LLM-as-a-Judge framework for evaluating text generation using checklists, 2024. URL <https://arxiv.org/abs/2403.18771>.
- [8] Tianci Liu, Ran Xu, Tony Yu, Ilgee Hong, Carl Yang, Tuo Zhao, and Haoyu Wang. OpenRubrics: Towards scalable synthetic rubric generation for reward modeling and LLM alignment, 2025. URL <https://arxiv.org/abs/2510.07743>.
- [9] Yang Liu, Dan Iter, Yichong Xu, Shuohang Wang, Ruochen Xu, and Chenguang Zhu. G-Eval: NLG evaluation using GPT-4 with better human alignment, 2023. URL <https://arxiv.org/abs/2303.16634>.
- [10] Yantao Liu, Zijun Yao, Rui Min, Yixin Cao, Lei Hou, and Juanzi Li. RM-Bench: Benchmarking reward models of language models with subtlety and style, 2024. URL <https://arxiv.org/abs/2410.16184>.
- [11] Anas Mahmoud, MohammadHossein Rezaei, Zihao Wang, Anisha Gunjal, Bing Liu, and Yunzhong He. Reward hacking in Rubric-Based reinforcement learning, 2026. URL <https://arxiv.org/abs/2605.12474>.
- [12] Arash Marioriyad, Omid Ghahroodi, Ehsaneddin Asgari, Mohammad Hossein Rohban, and Mahdieh Soleymani Baghshah. The judge who never admits: Hidden shortcuts in LLM-based evaluation, 2026. URL <https://arxiv.org/abs/2602.07996>.
- [13] Mohamad Amin Mohamadi, Tianhao Wang, and Zhiyuan Li. Honesty over accuracy: Trustworthy language models through reinforced hesitation, 2025. URL <https://arxiv.org/abs/2511.11500>.
- [14] Arjun Panickssery, Samuel R. Bowman, and Shi Feng. LLM evaluators recognize and favor their own generations, 2024. URL <https://arxiv.org/abs/2404.13076>.

- [15] José Pombal, Ricardo Rei, and André F. T. Martins. Self-Preference bias in Rubric-Based evaluation of large language models, 2026. URL <https://arxiv.org/abs/2604.06996>.
- [16] Maria Popp, Stefanie Reis, Selina Schießer, Renate Ilona Hausinger, Miriam Stegemann, Maria-Inti Metzendorf, Peter Kranke, Patrick Meybohm, Nicole Skoetz, and Stephanie Weibel. Ivermectin for preventing and treating COVID-19. *Cochrane Database of Systematic Reviews*, (6):CD015017, 2022. doi: 10.1002/14651858.CD015017.pub3. URL <https://pubmed.ncbi.nlm.nih.gov/35726131/>.
- [17] Pranav Rajpurkar, Robin Jia, and Percy Liang. Know what you don’t know: Unanswerable questions for SQuAD, 2018. URL <https://arxiv.org/abs/1806.03822>.
- [18] Delip Rao and Chris Callison-Burch. Autorubric: A unifying framework for Rubric-Based LLM evaluation on Non-Verifiable tasks, 2026. URL <https://arxiv.org/abs/2603.00077>.
- [19] Zihang Tian, Jingsen Zhang, Rui Li, Xiaohe Bo, Yuanzi Li, and Xu Chen. ARCO: Adaptive rubrics with Co-Evolution for Multi-Step LLM-Based agents, 2026. URL <https://arxiv.org/abs/2606.21262>.
- [20] Benedikt J. Wagner. Two axes of LLM abstention: Answer correctness and question answerability, 2026. URL <https://arxiv.org/abs/2607.08456>.
- [21] Beining Wang, Weihang Su, Hongtao Tian, Hao Kong, Tao Yang, Ting Yao, Qingyi Pan, Yueyue Wu, Qingyao Ai, Min Zhang, and Yiqun Liu. Co-Evolving LLM evaluators and policies via DynamicRubric, 2026. URL <https://arxiv.org/abs/2607.20083>.
- [22] Xiaohua Wang, Muzhao Tian, Yuqi Zeng, Zisu Huang, Jiakang Yuan, Bowen Chen, Jingwen Xu, Mingbo Zhou, Wenhao Liu, Muling Wu, Zhengkang Guo, Qi Qian, Yifei Wang, Feiran Zhang, Ruicheng Yin, Shihan Dou, Changze Lv, Tao Chen, Kaitao Song, Xu Tan, Tao Gui, Xiaoqing Zheng, and Xuanjing Huang. Reward hacking in the era of large models: Mechanisms, emergent misalignment, challenges, 2026. URL <https://arxiv.org/abs/2604.13602>.
- [23] Xuekang Wang, Zhuoyuan Hao, Shuo Hou, Hao Peng, Juanzi Li, and Xiaozhi Wang. Reproducing, analyzing, and detecting reward hacking in Rubric-Based reinforcement learning, 2026. URL <https://arxiv.org/abs/2606.04923>.
- [24] Koki Wataoka, Tsubasa Takahashi, and Ryokan Ri. Self-Preference bias in LLM-as-a-Judge, 2024. URL <https://arxiv.org/abs/2410.21819>.
- [25] Tianjun Wei, Wei Wen, Ruizhi Qiao, Xing Sun, and Jianghong Ma. RocketEval: Efficient automated LLM evaluation via grading checklist, 2025. URL <https://arxiv.org/abs/2503.05142>.
- [26] Zhepei Wei, Xiao Yang, Kai Sun, Jiaqi Wang, Rulin Shao, Jingxiang Chen, Mohammad Kachuee, Teja Gollapudi, Yiwei Liao, Nicolas Scheffer, Rakesh Wanga, Anuj Kumar, Yu Meng, Wen tau Yih, and Xin Luna Dong. TruthRL: Incentivizing truthful LLMs via reinforcement learning, 2025. URL <https://arxiv.org/abs/2509.25760>.
- [27] Minglai Yang, Xinyu Guo, Utkarsh Tyagi, Mian Zhang, Razvan Dumitru, Sunjie Hou, Yunzhong He, Daniel Yue Zhang, and Ying Liu. Rubric dropout: A simple way to mitigate reward hacking in Rubric-as-Reward RL, 2026. URL <https://arxiv.org/abs/2608.11669>.
- [28] Xianglin Yang, Bryan Hooi, Gelei Deng, Tianwei Zhang, and Jin Song Dong. Turning bias into bugs: Bandit-Guided style manipulation attacks on LLM judges, 2026. URL <https://arxiv.org/abs/2605.26156>.

- [29] Seonghyeon Ye, Doyoung Kim, Sungdong Kim, Hyeonbin Hwang, Seungone Kim, Yongrae Jo, James Thorne, Juho Kim, and Minjoon Seo. FLASK: Fine-grained language model evaluation based on alignment skill sets, 2023. URL <https://arxiv.org/abs/2307.10928>.
- [30] Skylar Zhai, Jingcheng Liang, and Dongyeop Kang. Abstain-R1: Calibrated abstention and Post-Refusal clarification via verifiable RL, 2026. URL <https://arxiv.org/abs/2604.17073>.
- [31] Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. Judging LLM-as-a-Judge with MT-Bench and chatbot arena, 2023. URL <https://arxiv.org/abs/2306.05685>.

A. Provenance and integrity contract

Three properties are **build-time hard invariants** enforced by `build_dataset.py --strict` (run in CI):

1. **Provenance.** `provenance_type` $\in$ {`real_web_grounded`, `hybrid`}; there are **zero pure-synthetic** environments. Any environment whose evidence contains a synthetic document is labeled `hybrid`. Counts: 158 `real_web_grounded` + 11 `hybrid`.
2. **Verifiability is derived, not asserted.** Each source carries a `verification_status` from a controlled vocabulary (`primary_fetched`, `mirror_fetched`, `secondary_corroborated`, `snippet_only`, `unverifiable`, `synthetic`); `source_verifiability` is computed deterministically (an environment is verifiable iff every evidence document has a `primary/mirror-fetched` reference, else `partially_verifiable`). Counts: 109 `verifiable` + 60 `partial`.
3. **Consistency.** `approved` $\Rightarrow$ `certificate_consistent`, and a valid gold answer may not contain a statistic absent from its packet. Both are hard errors.

All 169 environments are approved. The build emits `combined_environments.json[1]`, `controls.json[1]`, and `environments_all.json[1]`.

B. The gpt-5.6-sol resampling sub-study

The finding of this sub-study is that the 42% figure’s exploited set is draw-contingent even where its rate is stable, which is why §5 declines to extrapolate any $k = 1$ value. The rest of this section is the pre-registration and the record of what it did and did not establish. It holds the detail behind the two caveats attached to the 42% figure in §5:

The improvement is not a rubric-quality artifact. No 5.6 rubric meets the structural redraw bar that DeepSeek’s four malformed drafts met, none is cap-saturated (peak 16.5% of a 32k cap), and the nine shortest rubrics were exploited *more* often than the arm average, so terseness cuts against the 5.6 arms rather than for them. The improvement is concentrated almost entirely in `objective_conflict` (5–9/22 exploited, versus 14–16/22 for the older arms).

Why we do not know the sign of the resampling correction. Resampling moved the Claude arms downward (71→63, 96→82, 98→82), which invites the inference that 42% would also fall. All three of those arms were near the ceiling, however, where censoring alone predicts a downward move, so their common direction carries little information about an arm at 42%. This is the reason the main text forbids reading 42% against the resampled 63%.

Pre-registration and the subset. Under a pre-registration committed before any scoring agent ran, we drew two further gpt-5.6-sol rubrics per environment on a pre-specified stride subset of the cut, holding the chain fixed. Budget split the run into two pre-specified 10-environment blocks, the second fixed in writing before it was launched (amendment 2). The subset rate is near-flat across draws — 7/20, 7/20, 6/20 — but that is not per-environment stability: 3 of the 20 environments change verdict. The two blocks are reported separately as pre-committed, and they disagree — all three verdict-changing environments fall in the first block — though with three such environments in total that split is well inside chance and we offer no mechanism.

Why the 15% must not be read against §6.1’s 60%. The mixed-verdict fraction rises mechanically with the number of draws K : six draws give fifteen chances to disagree, three give three. The two figures are also measured at different n , on different generators, on different chains. At $n = 20$ this sub-study does not correct the 42%; it establishes that the *membership* of the exploited set is substantially draw-contingent, which a narrow interval on the rate would conceal. Per amendment 1 we report no confidence interval for this arm at any n .

The subset is where this arm ends. A pre-registered extension to 30 environments was authorized and launched, then halted at 6 of its 10 new environments on the second draw — one cell lost to a context-length failure, three blocked by a safety classifier on the harness’s subagent spawning, which we did not work around — and is recorded as **not executed** (amendment 4). The two predictions

amendment 3 had registered for it are therefore **unscored**, neither confirmed nor falsified. Its six completed cells are deliberately not folded in as $n = 26$: three of the four missing environments are `objective_conflict`, the type carrying most of the all-clean bucket, so dropping them would bias the rate *upward*. The final 15 environments were never authorized.

An exploratory observation that failed to replicate. Because the honest baseline answer is byte-identical across draws, any movement in *its* score is attributable to the rubric or to judge noise. Against a fixed-rubric repeat-scoring control (22 environments, byte-identical rubrics, same chain and period), the rubric draw moved it $2.8\text{--}5.0\times$ more than chain noise at $n = 10$, with one of three draw pairs at $p = 0.025$. Doubling to $n = 20$ halved the effect to $1.8\text{--}2.2\times$ and lifted every pair above $p = 0.05$. The surviving claim is a direction-consistent but **unconfirmed** $\approx 2\times$ effect — the ordinary fate of an estimate first measured at the moment it was discovered, which is why the analysis is labelled post hoc and excluded from the pre-registration.

C. Open-weights parity: per-type profile and generation integrity

Two structural observations reinforce that the DeepSeek result of §5 is parity rather than noise in a single number. DeepSeek’s **per-type failure profile is within one environment of Opus’s in every impossibility type**, and DeepSeek and GPT-5.5 agree on **32/45** environments (24 exploited by both, 8 by neither). Hard-set difficulty is therefore largely a property of the *environment*, not of the generator — which is what one wants from a benchmark, and which makes the frontier-vs-mid gap (69–71% vs 96–98%) harder to attribute to any one vendor’s training.

We also flag one asymmetry that is *not* about rubric quality: on first draws the open model emitted **4 structurally malformed rubrics / 45** (one empty, two silently truncated, one duplicated) against GPT-5.5’s **0 / 169**, every one of them reporting a normal `finish_reason`. These were regenerated at identical settings under a purely structural criterion, fixed and logged before the redraws and applied without reference to any score.

D. Counts behind the headline claims

On (a), room for improvement. The certificate-faithful reference yields 0/45 exploits under the tested chain, while generated rubrics retain positive rates. This empirical gap motivates further evaluation; it is not a verified optimum or a property of rubric generation independent of the attacker and Oracle (§I, §M).

On (b), the transitive chain. DeepSeek 31/45 $\approx$ GPT-5.5 30/45 (paired McNemar, $p = 1.00$), and GPT-5.5 30/45 sits two environments under Opus 32/45. That second link is **unpaired**: the Opus Hard-45 arm is one of the three $k = 1$ Claude runs with no surviving per-cell record (see the reproduction appendix), so no McNemar is available for it and we compare counts only. The three quantities differ by at most 2 environments, which is within the 13.6% repeat noise, so the conclusion holds — as a chain of two weak links, which is why we judged that boundary indeterminate rather than resolved.

Against the contemporaneous 5.6 arms the same open-weights generator sits 12 environments below on the hard set, and the one 5.6 arm we also measured on the unbiased cut sits 10 below it there (sol 15/150 vs DeepSeek 25/150, $p = 0.0129$). That is why (b) is stated as level-specific parity rather than as an open-weights ranking claim.

On the multiplicity discipline behind the Full-150 block. The nominal paired McNemar values against DeepSeek’s 25 are sol $p = 0.0129$ (discordant 2 vs 12), luna $p = 0.0129$ (2 vs 12) and terra $p = 0.0352$ (3 vs 12), all unadjusted; the three within-5.6 comparisons are discordant 4 vs 5, 4 vs 4 and 5 vs 4, every $p = 1.0$. Under Holm at $\alpha = 0.05$ across that family of six the smallest p must clear $0.05/6 = 0.0083$ and 0.0129 does not, so the step-down stops at the first step. Completing luna weakened the corrected claim while strengthening the raw evidence, and we report it that way. With three arms that family had three comparisons and sol vs DeepSeek cleared $0.05/3 = 0.0167$; luna’s arm adds a second independent $p = 0.0129$ against DeepSeek and three null within-5.6 comparisons,




Figure 4: Rubric-resampling confidence intervals on HARD-45. Dots are resampled point estimates with cluster-bootstrap 95% CIs over environments; $\times$ marks the single-draw ($k=1$) value the first leaderboard reported. For these three generators the resampled means are lower than the original draws, and Sonnet/Haiku CIs overlap. This does not establish a universal bias direction or equivalence. All three generated-rubric intervals sit far above the certificate-faithful reference (0/45 under the tested chain); Sonnet and Haiku also sit above the naive decisiveness proxy (64%), while the Opus interval straddles it.

and the multiplicity burden of the latter is what removes the surviving result. Restricting the family post hoc to the three 5.6-vs-DeepSeek tests would restore it, and we do not do that: the family is the set of comparisons the block reports, it was defined before luna’s arm existed, and narrowing it after seeing which definition preserves a significance verdict is exactly the choice the correction is meant to bind.

On the first conjunct. Across the twenty-one arms with surviving per-cell records (2,100 environment $\times$ arm cells), 522 cells violate the certificate and 514 are judged exploited. Thirteen cells separate the two: eight where the adversarial answer violated but scored *below* the honest baseline, and five decided by the $\geq$ tie-break alone — an exact tie between the adversarial and honest answers, resolved as an exploit by condition (a). All rates in the paper use the pre-registered $\geq$ rule. Under a strict $>$ rule, exactly one displayed equality in §5 breaks and no separability verdict changes. The closest Full-150 pair, GPT-5.6-terra vs. sol, goes from 16–15 to a tie — already reported as not separable ($p = 1.0000$). Sonnet 5’s Hard-45 arm goes from 30/45 to 29/45 while GPT-5.5 and DeepSeek have no tie cells and do not move, so the reported GPT-5.5 67% = Sonnet 5 67% becomes 67% vs. 64%; that pairing does not separate under the strict rule either ($p = 0.7539$), and Sonnet 5 stays short of all three GPT-5.6 arms either way. One pairing crosses $\alpha = 0.05$ under the strict rule — Sonnet 5 vs. gpt-5.4-mini, $p = 0.0654 \rightarrow 0.0386$ — widening a gap already reported in that direction rather than reversing one.

E. Judge robustness

Holding a frozen GPT-generated rubric, adversarial answer, and Oracle verdict fixed, we varied only the judge across Haiku, Sonnet, and Opus on 60 chains (Table 5). The score comparison $J(a^*; r) \geq J(a_0; r)$ changes on 4/60 chains, but all four have non-violating Oracle labels, so the final exploit verdict is unchanged on 60/60. This checks judge sensitivity on the sampled chains; it does not establish ranking invariance across all generators.

Only 10/60 chains are Oracle-positive, limiting the opportunities for judge changes to affect the conjunction. Across the 2,100 leaderboard cells with surviving records, eight Oracle-positive attacks score below the baseline and five exploit verdicts depend on an exact tie (§D). Verification labels therefore carry substantial weight in the reported rates. The 40-item human calibration observes one false positive and one false negative (§F); it cannot rule out either error direction. The frozen-answer Oracle swap in §M directly measures sensitivity to the verification configuration.

Table 5: Robustness and calibration summary. Each row perturbs exactly one element of the chain and reports what survives.

analysis	perturbation	result
Judge (§E)	rubric, answer, and Oracle verdict fixed; judge $\in$ {Haiku, Sonnet, Opus}	score comparison changes on 4/60 chains, exploit verdict on 0/60; only 10 Oracle-positive chains limit power
Rubric draw (§6.1)	$k=3$ rubrics/env ($k=6$ Opus), cluster bootstrap $B=10^4$	Opus 63.3% [53.7, 73.0] < Sonnet 81.5% [70.4, 91.1] $\approx$ Haiku 82.2% [73.3, 90.4]; Opus–Sonnet -18.1% [$-27.8, -8.5$]; Sonnet $\approx$ Haiku not separable
Human calibration (§F)	40-item blind human rating vs. the Opus Oracle	agreement 38/40, $\kappa = 0.89$; overall specificity 25/26; predefined honest subgroup 13/13; one false positive and one false negative
Attacker (§G)	attacker Opus $\rightarrow$ GPT-5.5, everything else fixed (paired)	gap survives: Sonnet–Opus $+18\%$ [$+2, +36$], Haiku–Opus $+22\%$ [$+4, +40$]; per-env attacker agreement 80%
Composition (§H)	per-type rates on FULL-150 + type-stratified bootstrap	ranking holds within 5/6 types; composition-neutral gap modest (Haiku–Opus $+5\%$ [$+1, +9$]); discrimination concentrated in evidence_insufficient + objective_conflict
Calibration (§I)	generated rubric $\rightarrow$ certificate-faithful and naive-proxy references	certificate-faithful 0/45; naive proxy 64% [49, 78]; seven generated arms exceed the proxy under the tested chain; shared certificates limit independent validation
Optimization pressure (§J)	instructed adversary $\rightarrow$ best-of-N over a neutral policy ($N=5$, no adversary)	rubric-argmax misses an available honest answer 0% observed (13 recoverable envs) and beats random selection ($\Delta = -16.7\%$ [$-30.0, -5.6$]); 28% of envs yield no honest sample at all
Generation prompt (§6.2)	neutral prompt $\rightarrow$ + one generic “penalize overclaiming” sentence; 3 arms, all generated over raw HTTP	exploited Opus 5 36% $\rightarrow$ 22%, GPT-5.6-sol 42% $\rightarrow$ 36%, Sonnet 5 67% $\rightarrow$ 49%; paired $-13.3/-6.7/-17.8$ pts, McNemar $p = 0.109/0.508/0.039$ (unadjusted); all three residual rates exceed the 0/45 certificate-faithful reference
Re-scoring drift (§5)	same generator scored twice, in each of the two scoring rounds (Δ on envs scored both times, $n \geq 40$)	gpt-5.5 32/45 $\rightarrow$ 30/45 ($\Delta = -2$, $p = 0.6875$), gpt-5.4-mini 40/45 $\rightarrow$ 37/45 ($\Delta = -3$, $p = 0.2500$); pre-registered verdict indeterminate; both move <i>down</i> , opposite a stronger-attacker account; same-period repeat noise 3/22 = 13.6%. HARD-45 only — no FULL-150 bridge arm exists




Figure 5: Rubric-draw variance per environment on HARD-45. Each environment is classified by how many of its k rubric draws were exploited. The *mixed* band is where a single-rubric measurement is a coin flip: for Opus it covers 27 of 45 environments, so which rubric one happens to draw decides the verdict on 60% of the hard set. This is why the leaderboard reports resampled rates with intervals rather than single draws.

F. Oracle human calibration

A benchmark author rated a 40-item, Oracle-blind sheet containing honest and adversarial answers across all six types, shuffled under opaque codes. The sheet was stratified by the original Oracle labels (14 positive, 26 negative); the rater did not see those labels, the answer type, or the certificate checklist. Relative to this rater, the confusion matrix is 13 true positives, 25 true negatives, one false positive (CAL-07), and one false negative (CAL-14). Agreement is 38/40 (95%), Cohen’s $\kappa = 0.89$, sensitivity is 13/14, and **overall specificity is 25/26 (96.2%)**.

All 13 predefined honest-baseline items agree with the rater, whereas agreement in the adversarial subgroup is 25/27. The 13/13 subgroup result is not overall specificity: one adversarially authored answer was judged honest by the rater but violating by the Oracle. These observations support agreement on this small sample, not the absence of false positives or systematic bias. There is one rater, so inter-annotator agreement is unmeasured. Labels apply to the exact rated text; they cannot validate a different attack answer for the same environment in a later experiment.

G. Held-out attacker (hard-set circularity)

The hard set is *selected*, and the leaderboard *rank*s, using the fixed Opus attacker — a potential circularity. (The judge axis is covered by §E.) We tested the attacker axis with a paired design: per (environment, generator) we drew one neutral rubric, then attacked it with **both** Opus (control) and a held-out cross-vendor attacker, **GPT-5.5**, with the same judge and oracle on both — so only the attacker changes. On the 45 hard environments, the **generator gap survives the held-out attacker**: Sonnet – Opus = +18% [+2, +36] and Haiku – Opus = +22% [+4, +40] under GPT-5.5 (both CIs exclude 0). Per-environment, the two attackers reach the same exploited verdict on the identical rubric **80%** of the time. The tested generator gaps therefore persist with this held-out attacker, but this does not rule out sensitivity to other attackers or remove the hard-set selection effect. (This arm is single-rubric; the gap is also established at $k \geq 3$ in §6.1.)

H. Composition / per-type analysis

The hard set is objective_conflict-heavy (22/45). To test whether the ranking is a composition artifact, we computed per-type exploit rates on the **unbiased Full-150 cut** (every environment under all three Claude generators) and a type-stratified bootstrap that equal-weights types.

TYPE	N	OPUS	SONNET	HAIKU
evidence_insufficient	25	52%	60%	60%
objective_conflict	25	16%	24%	24%
exhaustiveness_impossible	25	0%	16%	12%
missing_support_citation	25	4%	8%	8%
counterevidence_trap	25	8%	0%	4%
false_premise	25	0%	0%	0%

The ranking direction ($\text{Opus} \leq \text{Sonnet}$, Haiku) holds **within 5 of 6 types**; the lone exception (counterevidence_trap) is a single environment (2/25 vs. 0/25), i.e., noise. The **composition-neutral gap is modest**: equal-weighting types, Haiku – Opus = +5% [+1, +9] (excludes 0) and Sonnet – Opus = +5% [+0, +9] (touches 0). This confirms the 18–20-point hard-set gaps are **selection-amplified** versions of the same direction, not the unbiased effect size. Finally, discriminating power is **concentrated**: only evidence_insufficient (52–60%) and objective_conflict (16–24%) substantially stress rubrics, while false_premise is 0% across all generators (non-discriminating under this chain). The objective_conflict over-representation is thus downstream of a real fact — real-grounded impossibility reliably stresses rubrics in only two types — and is a presentation matter (pair hard-set numbers with the per-type table), not a validity threat. No rebalancing improves discrimination, since the under-represented types barely exploit.

I. Certificate-faithful and naive-proxy calibration

To anchor the absolute exploit rate we ran the same chain under two *fixed*, non-generated rubrics on the hard set. The **certificate-faithful reference** is each environment’s certificate rendered as a rubric (reward required_honest_behavior and honest framing; automatic-zero on any unsupported claim or reward hack). The **naive proxy baseline** is a generic rubric (“we want a decisive, well-cited answer; penalize hedging”), identical across environments. Results under the fixed chain are **certificate-faithful 0/45** and **naive proxy 64% [49, 78]**. Zero observed exploits is not evidence of zero population risk. The certificate-derived rubric provides an empirical reference showing that these tasks can receive more faithful rewards under the tested chain. Because rubric construction and Oracle verification share the same certificate, this comparison is not an independent validation of certificate correctness or an Oracle-invariant optimum. A secondary, more surprising observation: on the single-draw board **seven of the eleven generated arms are exploited at least as often as the naive proxy**, not less (both are $k = 1$ values) — plausibly because a generated rubric’s task-specific criteria (“state the threshold,” “report the score”) act as an **attack roadmap** that forces the violating claim, whereas a generic rubric gives the attacker no such target. Specificity, which makes a rubric look better, is what makes it exploitable. **Opus 5 and the three GPT-5.6 arms are the exception** (36%, 42%, 42%, 51%): they fall below the proxy, so the roadmap effect is not a necessary property of task-specific rubrics. §B rules out length and form as the difference; what does explain it we have not identified. (*This direction held across two naive-proxy designs; magnitude is single-rubric.*)

J. Realistic selection pressure: best-of-N without an adversary

Every exploit number so far comes from an attacker *instructed* to ignore honesty — worst-case exploitability. Rubric-as-reward systems, however, first meet a rubric as a **selection reward over a policy’s own samples**. We therefore removed the adversary entirely: for 18 type-stratified Hard-45 environments, the strongest generator (Opus) wrote one neutral rubric; a fixed neutral policy (Sonnet) that never sees the rubric drew $N = 5$ answers; the fixed judge scored each sample independently under the generated rubric and under the certificate-faithful rubric; the fixed oracle checked every sample for certificate violation.

Three results. (i) **The policy itself is the dominant realistic failure**: in 5 of 18 environments (28%), *none* of the N samples was honest — no selection rule can recover honesty that the policy never produces (the answerer-side counterpart of AbstentionBench [5]’s finding). (ii) **Where an honest sample exists, the generated rubric finds it**: on the 13 recoverable environments the rubric-argmax selected a certificate-violating answer 0/13 times (tie-breaking bracket 0–8%), matching the certificate-faithful rubric on the same samples. (iii) **The generated rubric beats no reward at all**: best-of- N violation 27.8% vs. 44.4% for random choice, $\Delta = -16.7\%$ [−30.0, −5.6], CI excluding 0 — used as a selection reward under benign sampling, the rubric actively steers *toward* honesty.

This small best-of- N study checks reward selection under one neutral sampling policy. It does not establish that exploitation requires an adversary, that stronger search increases failures monotonically,

or that the adversarial benchmark is a ceiling for RL training. Zero failures among the 13 recoverable environments is a sample result, not a zero-risk guarantee. The study covers one generator, one rubric per environment, $N = 5$, and one policy; larger-sample selection and end-to-end training remain unmeasured.

K. Negative control: is the exploit rate bought with over-refusal?

A generator could in principle score 0% exploited by writing a degenerate rubric that rewards hedging unconditionally. We check that it does not. On the 48 answerable controls, three Claude generators (Opus 4.8, Sonnet 4.6, Haiku 4.5) write a rubric under the **same** neutral prompt used on the impossible side — the generator is not told the task is answerable — and the **same** fixed literal judge scores the shipped `correct_grounded` answer against two shipped distractors, a polished `over_refusal` and a `wrong_confident`. There is no attacker and no oracle: control environments carry no certificate, so the two hardest stages of the chain (§4.2) are absent by construction, and the metrics are single conditions rather than the conjunction that defines exploitation. Result: **false-refusal 0%, wrong-confident 0%, correct strictly first in 144/144 cells**.

What this does not establish. The same controls also return **0/36** on both metrics for the *naïve proxy* rubric — the generic “be decisive, penalize hedging” rubric that §I measures at **64% exploited** on the hard impossible cut. A check that passes the paper’s own designated bad rubric does not separate rubric quality at this adversary strength; it separates gross degeneracy from everything else. We therefore report it as a floor, not as evidence that the leaderboard ordering reflects discrimination rather than caution, and we do not compare the 0% here to the 63–82% on the impossible cut as if the two were measured at comparable pressure.

Why the distractors are fixed rather than adaptive. This was not a shortcut; an adaptive version was built and does not run. The pilot chain (`leaderboard_neutral_run.js`) carries an attacker branch for answerable environments — write a polished over-refusal designed to score high under this rubric — and we ran it on 12 of the 48 controls $\times$ the same three generators. On **11 of 36 cells the attacker declined the instruction** and returned the correct grounded answer instead (“declined to fabricate a false ‘evidence insufficient’ refusal for a clearly answerable question”), which the $\geq$ rule then records as a false-refusal hit: a metric artifact, not a rubric failure. On the 25 cells where the attack actually executed, false-refusal was **0/25**. The obstruction is structural rather than incidental — on an answerable question with decisive evidence, the over-refusal attack is the request to write persuasive evasion about a settled fact, which a frontier attacker refuses roughly a third of the time. A stronger over-refusal adversary would need an attacker that is not safety-trained against exactly this, and we do not have one; the fixed distractors are what make the axis runnable at all. This is a reusable observation for anyone building a false-refusal control: the impossible and answerable directions are not symmetric, because only one of them can be attacked by asking a capable model to argue for it.

L. Self-play and a held-out attacker

The Opus 5 generator arm uses the same model family for rubric generation, attack, and verification. Its original Hard-45 rate is 16/45 (36%), with no score ties or below-baseline violating attacks. Its exploit set is a subset of the Sonnet 5 and DeepSeek sets in this run (discordant pairs 0 versus 14 and 0 versus 15; exact McNemar $p = 0.0001$ for both). It does not separate from the three GPT-5.6 arms ($p = 0.55, 0.58, 0.09$).

Three possible artifacts tested. First, the attacks reach a mean judge score of 99/100, have a median length of 8.5k characters, and contain no refusals. High reward and long answers show engagement with the rubric, but do not establish optimal attack strength. Second, using the same Opus attacker across generators holds the author’s model family fixed; it does not preclude Oracle bias that interacts with answer content or rubric-induced style. Third, none of 45 attacker transcripts verbalizes self-recognition; 17 mention scoring caps and seven describe following the rubric over the

instruction to be unequivocal. These reasons are descriptive evidence and cannot exclude unspoken authorship effects.

Matched attacker swap. Holding the rubrics, Haiku judge, Opus Oracle, and certificates fixed, swapping the attacker to GPT-5.5 changes Hard-45 exploitation from 16/45 to 15/45 (−2.2 percentage points, paired 95% CI [−13.3, +8.9]; discordant pairs 4 versus 3, exact $p = 1.00$). Itemwise agreement is 38/45. On Full-150 the rate changes from 12/150 to 10/150 (−1.3 points [−4.7, +2.0]; discordant pairs 4 versus 2, $p = 0.6875$), with agreement 144/150. The cuts overlap: the 26 hard environments inside Full-150 carry every positive label, while both attackers yield 0/124 on its remaining environments. These are overlapping observations rather than independent replications.

The tested attacker swap yields similar aggregate counts, but nonsignificance does not establish equivalence or rule out stronger attackers. Both arms retain the Opus Oracle, leaving verification sensitivity unresolved. Section M addresses that axis using the frozen GPT-5.5 answers: its historical comparison is 15/45, not the 16/45 obtained with Opus-authored answers. The leaderboard remains conditional on its original evaluation chain; this experiment does not establish ranking invariance under another Oracle.

M. Sensitivity to the Oracle configuration

Frozen design. We retain the 45 Hard-45 environments, Opus 5 rubrics, evidence packets, certificates, GPT-5.5 attack answers, reference baselines, and Haiku scores from the held-out-attacker arm of §L. Each Oracle evaluates all 90 texts independently. The reference baselines and attacks are presented as opaque shuffled items; the Oracle receives the question, evidence, certificate fields, and candidate answer, without role labels, the rubric, judge scores, or previous verdicts. The prompt explicitly introduces the certificate requirements with “an honest answer MUST:”. For environment i and Oracle m , exploitation is $E_{im} = T_i \wedge V_{im}$, where $T_i = [J(a_i^*; r_i) \geq J(a_{0i}; r_i)]$ is frozen and V_{im} is the new attack-violation label. The score gate is true on 44/45 environments. None of the Oracles flags the remaining attack, so attack-violation and exploit counts coincide here. Baseline rejection is reported separately; it does not remove environments or change the primary denominator.

Oracle configurations and execution. O-A uses anthropic/claude-opus-5, O-B openai/gpt-5.6-sol, and O-C google/gemini-3.8-flash, with provider routes pinned to Anthropic, OpenAI, and Google AI Studio. Each call is stateless, uses a single user message without tools, requests the same two-field JSON response, and allows 16,000 output tokens. Temperature, top- p , seed, and reasoning settings are omitted and inherit model defaults. The configuration was frozen before the formal run on 15 September 2026. All 270 formal judgments succeeded on their first attempt with normal completion; three smoke calls bring the request total to 273. Saved usage records sum to \$2.99. Protocol, input hashes, per-call records, and analysis outputs are retained with the study artifacts.

Oracle	Exploitation	Wilson 95% CI	Baselines rejected
O-A: Claude Opus 5	15/45 (33.3%)	[21.4, 47.9]%	0/45
O-B: GPT-5.6-sol	34/45 (75.6%)	[61.3, 85.8]%	3/45
O-C: Gemini-3.8-flash	30/45 (66.7%)	[52.1, 78.6]%	0/45

Paired differences. O-B exceeds O-A by 42.2 percentage points (paired bootstrap 95% CI [28.9, 57.8]); O-C exceeds O-A by 33.3 points [20.0, 46.7]. Two-sided exact McNemar tests, Holm-adjusted across all three Oracle pairs, give $p = 1.14 \times 10^{-5}$ and $p = 1.22 \times 10^{-4}$. O-B minus O-C is 8.9 points [−2.2, 22.2], $p = 0.289$; this is not evidence of equivalence. Intervals use 10,000 paired environment resamples with seed 20260915. All 15 O-A-positive attacks are also positive under O-B and O-C in this run. The latter two sets are not nested: six attacks are positive only under O-B and two only under O-C. Unanimous, majority, and any-Oracle rules yield 15/45, 28/45, and 36/45, respectively. These describe alternative decision rules, not validated truth labels or bounds on the true rate.

Comparison with the historical Oracle. Fresh O-A judgments match the historical 15/45 aggregate but agree itemwise on 41/45 attacks, with two flips in each direction. The historical evaluation used an Opus harness agent with a system prompt and an unpinned model alias; O-A uses the pinned API configuration above. Thus this is an aggregate match across implementations, not an exact replication. The two minimal-format attacks, “3” and “Yes”, each score 100 versus a baseline score of 10 and are flagged by all three Oracles. Their task requests explicitly require a numeric cell or a decisive word, so brevity alone does not invalidate them. Excluding both descriptively gives 13/43, 32/43, and 28/43; the paired discordances and exact tests are unchanged. The original attack logs lack completion metadata, so they cannot independently establish whether the generation was truncated. Primary results retain all 45 answers.

Baseline and certificate consistency. O-B rejects three reference baselines accepted by O-A and O-C. In `evid_insuff_m004`, the baseline omits certificate-required cautions about a small, single-organization study. In `stats_objective_conflict_016`, it attributes a roughly 22% figure to UBA rather than CEW and adds an electricity-metric interpretation not established in the packet. In `stats_objective_conflict_019`, it omits the required first ACGR value (79% for 2010–11) and the AFGR estimation-method distinction. These cases expose inconsistencies between the frozen baselines and the written requirements; the rejection counts are therefore not estimates of Oracle false-positive rates. Inputs and labels remain unchanged, so this audit does not retroactively redefine the benchmark.

Interpretation and limits. Some reasons emphasize overall honesty while others emphasize unmet requirements, source attribution, or unsupported numerical extrapolation. Because the prompt already says MUST, the disagreement cannot simply be attributed to an absence of mandatory language. Reasons suggest possible interpretations but do not identify what causes the rate differences. All three Oracles read the same certificates, leaving shared certificate errors invisible to agreement; no new independent human labels adjudicate these answers. Earlier calibration labels apply only to their exact texts, not to a different attack sharing a task ID. O-B shares a vendor with the attacker, O-C is a Flash model, and model defaults and historical interfaces differ, so vendor, model capability, and implementation effects are not separately identified. This is one frozen generator/attacker condition with one judgment per text and Oracle. It quantifies sensitivity of the measured rate, without establishing the most accurate Oracle, a true exploitation rate, repeat-run stability, or the ordering of the eleven generators under alternative verification.